

Page 01 — Homepage / (index.html)

Audited: 3 September 2026 · live site + repository `bccody03/blog-site` @ `claude/blakecody-com-access-yi26f9`

Live URL: `https://blakecody.com/` — HTTP 200, GitHub Pages via Fastly, `Last-Modified: 2026-07-20`

Method: static HTML parse · headless Chromium render (Playwright) · live HTTP header + redirect probes · DNS · full read of `app.js`

Status: Audit only — **no code changes have been applied**. Every entry below is a proposed change with the exact edit.

32

CHANGES PROPOSED

11

HIGH

14

MEDIUM

7

LOW / INFO

Scorecard

Area	Grade	Headline
Technical	C	Clean, valid markup and correct canonical/robots/sitemap basics — undermined by HTTP not redirecting to HTTPS, all real content rendering client-side from a third-party API, and zero structured data.
Content / text	D	146 words of indexable static copy. The H1 and title spend the page's strongest signals on a brand term you already own, and 100% of article content lives on Substack, not on this domain.
Tracking	D	GoatCounter counts pageviews and nothing else. Not one conversion on the site — Substack clicks, App Store taps, chapter downloads, reflections — is measured.
Performance	C	Small HTML, gzip on, sensible asset count. But the LCP image is JS-injected, the font CSS is a render-blocking third-party dependency, and a 6-second intro animation owns the first visit.

Do these in this order

P0 — this week (≈2 hours, biggest return)

1. **T1** Turn on Enforce HTTPS in GitHub Pages settings — *2 minutes*
2. **T3** Delete the fake sample-post fallback — *15 minutes*
3. **K1** Wire GoatCounter events for the five conversions — *30 minutes*
4. **C1 C2 C6** Rewrite title, H1 and meta description — *20 minutes*
5. **T5** Add the JSON-LD `@graph` block — *30 minutes*

P1 — this month

1. **C3** Write 120–180 words of real body copy onto the homepage
2. **T4** Self-host the fonts (or make the font CSS non-blocking)
3. **T6** Move the hero background into CSS and preload it
4. **T10** Fix image weights and add width/height
5. **K3** UTM-tag every outbound Substack and App Store link
6. **T8 C5 C8** Internal links to `/`, post titles to `h3`, real alt text

P2 — the strategic one (do it once, changes everything)

1. **C4 + T2** Host the essays on blakecody.com and prerender them at build time. Until this happens the domain has no article content to rank, and every essay you publish builds Substack's authority instead of yours.

A. Technical

T1 HTTP is served directly — no redirect to HTTPS HIGH

NOW

```
$ curl -I http://blakecody.com
HTTP/1.1 200 OK          <- serves the page. No 301. No Location header.
server: GitHub.com
```

Google treats `http://` and `https://` as two separate URL sets, so the site currently exists twice. There is no HSTS, anyone landing on an `http://` link sees Chrome's "Not secure", and the GoatCounter script (see **K2**) is loaded protocol-relative, so on those loads your analytics go over plaintext too.

CHANGE TO

GitHub repo → **Settings** → **Pages** → **Enforce HTTPS** (checkbox). The certificate is already provisioned, so this is a two-minute change with no code.

VERIFY

```
$ curl -I http://blakecody.com
HTTP/1.1 301 Moved Permanently
location: https://blakecody.com/
```

T2 The only real content on the page is rendered client-side from a third-party API HIGH

NOW

The served HTML contains an empty container and a loading string:

```
<div class="state" id="state">Loading the latest...</div>
<div class="post-list" id="post-list" data-limit="3"></div>
```

Posts are injected by `app.js` → `init()` → `loadFromSubstack()`, which `fetches` `api.rss2json.com` (a public relay, with the API key hardcoded in the client-side file) and only then builds the cards.

Measured: **146 words** of static indexable text vs **242 words** after render.

Googlebot does render JavaScript, but on a deferred second pass and within a budget — and a cross-origin API call inside that pass makes your freshest content non-deterministic to the crawler. Bing, and essentially every AI/LLM crawler, do not execute JS at all: they see "Loading the latest..." and nothing else.

CHANGE TO

1. **Prerender at build time.** A GitHub Action on a cron (say every 6 hours) fetches the Substack RSS, writes the latest 3 cards into `index.html` and 12 into `articles.html`, and commits. The content is then in the served HTML for every crawler, on a static host, at zero runtime cost. `app.js` keeps its fetch only to refresh what is already there.
2. **Interim (10 minutes):** add a `<noscript>` block with the three most recent post titles and links, plus the `ItemList` JSON-LD from **T5**.

Note: the hardcoded `rss2json` key is public in `app.js` and served to every visitor. That is inherent to this pattern, but it means anyone can spend your quota — another reason to move the fetch to build time, where the key can live in GitHub Actions secrets.

T3 When the feed fails, the homepage publishes three articles that do not exist HIGH

NOW

Headless render of the live markup, feed unreachable:

```
state:      "Couldn't reach the Substack feed right now. Showing sample posts."
postCards: 3
  JUNE 10, 2026  "Why I Started Writing in Public"      -> href="#"
  MAY 22, 2026   "The Case for Doing Less, Better"     -> href="#"
  MAY 1, 2026    "Notes on Building Small Things"      -> href="#"
```

These come from the `SAMPLE_POSTS` array in `app.js`. They are invented titles with invented dates and dead `#` links, and they are shown to real visitors and to Googlebot whenever `rss2json` is down, rate-limited, or the key lapses. If a render happens during that window, fabricated articles get indexed under your name — a credibility problem well beyond SEO.

CHANGE TO

```
} catch (err) {
  console.error(err);
  track('feed-error'); // see K1 – get alerted instead of surprised
  setState('Latest posts are on Substack →', false);
  // render nothing. Never render placeholder articles.
}
```

Keep `SAMPLE_POSTS` if you want it for local development, but gate it: `if (new URLSearchParams(location.search).has('demo'))`.

Could not verify the live feed state from this environment — `api.rss2json.com` is blocked by the sandbox's egress policy. Open `blakecody.com` in a browser and look for the "Showing sample posts" line under "Latest articles".

T4 Render-blocking third-party font CSS sits on the critical path

HIGH

NOW

```
<link href="https://fonts.googleapis.com/css2?
family=Fraunces:ital,opsz,wght@0,9..144,400;0,9..144,600;
1,9..144,400&family=Inter:wght@400;500;600&display=swap" rel="stylesheet" />
```

Two families, six faces including an italic, fetched from a third party, render-blocking, in the head. In the headless render where that host was unreachable, **first-contentful-paint was 12.9 seconds** — the page painted nothing at all until the request gave up. In normal conditions it is fast, but it is still a single point of failure that can hold your entire first paint hostage, and it costs a DNS lookup + connection + round trip before anything appears.

CHANGE TO

Best: self-host. Download the woff2 files into `/fonts/`, declare them with `@font-face { font-display: swap }` in `styles.css`, and preload the two faces used above the fold:

```
<link rel="preload" href="/fonts/fraunces-600.woff2" as="font" type="font/woff2" crossorigin>
<link rel="preload" href="/fonts/inter-400.woff2" as="font" type="font/woff2" crossorigin>
```

Minimum: keep Google Fonts but make it non-blocking, and drop the faces you do not use (the Fraunces italic and at least one Inter weight):

```
<link rel="preload" as="style" href="https://fonts.googleapis.com/css2?... "
onload="this.onload=null;this.rel='stylesheet'">
```

`display=swap` is already set — that part is correct.

T5 No structured data anywhere on the page HIGH

NOW

`document.querySelectorAll('script[type="application/ld+json"]').length === 0`. There is no machine-readable statement that Blake Cody is a person, that *Aligned* is a book he wrote, or that My JRNL is his app. Nothing for a Knowledge Panel, nothing for sitelinks, nothing for an answer engine to attribute a claim to.

CHANGE TO

One block before `</head>`. Fill in the `sameAs` list with every profile you actually control:

```
<script type="application/ld+json">
{
  "@context": "https://schema.org",
  "@graph": [
    { "@type": "Person", "@id": "https://blakecody.com/#person",
      "name": "Blake Cody", "url": "https://blakecody.com/",
      "jobTitle": "Author", "image": "https://blakecody.com/logo.png",
      "knowsAbout": ["discipline", "intentional living", "journaling", "personal growth"],
      "sameAs": ["https://bccody.substack.com",
        "https://apps.apple.com/us/app/my-jrnl-journal-habits/id6788176326"] },
    { "@type": "WebSite", "@id": "https://blakecody.com/#website",
      "url": "https://blakecody.com/", "name": "Blake Cody",
      "inLanguage": "en-US", "publisher": { "@id": "https://blakecody.com/#person" } },
    { "@type": "Book", "@id": "https://blakecody.com/#aligned",
      "name": "Aligned: Lead with Love, Live with Intention",
      "author": { "@id": "https://blakecody.com/#person" },
      "url": "https://blakecody.com/book.html" },
    { "@type": "WebPage", "@id": "https://blakecody.com/#webpage",
      "url": "https://blakecody.com/", "isPartOf": { "@id": "https://blakecody.com/
#website" },
      "about": { "@id": "https://blakecody.com/#person" },
      "primaryImageOfPage": "https://blakecody.com/og-image.png" }
  ]
}
</script>
```

Add a `SoftwareApplication` node for My JRNL on whichever page promotes it hardest, and an `ItemList` of the latest posts once T2 puts them in the HTML. Validate at `validator.schema.org` and in Search Console's Rich Results Test.

T6 The LCP image is invisible to the browser's preload scanner MEDIUM

NOW

The full-screen hero (`hero-bg.jpg` , 1600×1000, 90KB) is not referenced in the HTML or the CSS. `app.js` sets it as an inline style after it parses:

```
function setCover(url) {  
  if (url && els.coverImg) els.coverImg.style.backgroundImage = `url("${url}")`;  
}
```

So the largest element on the page cannot begin downloading until HTML parse → JS download → JS execute have all finished. That is a wholly avoidable delay on your Largest Contentful Paint.

CHANGE TO

`CONFIG.coverImage` is a constant (`"hero-bg.jpg"`), so declare it statically and let JS only override it in the "use the latest post's cover" branch:

```
/* styles.css */  
.cover-img { background-image: url("hero-bg.jpg"); }  
  
<!-- index.html, in <head> -->  
<link rel="preload" as="image" href="hero-bg.jpg" fetchpriority="high">
```

T7 A six-second intro animation owns the first visit MEDIUM

NOW

`buildIntro()` runs a 4.6-second choreography and the overlay is removed at `setTimeout(..., 6000)`. A full-screen element covers the page for ~6s on the first load of every browsing session — which is exactly the load a visitor arriving from Google gets.

Credit where due: it is `sessionStorage`-gated so it does not replay, it respects `prefers-reduced-motion`, and it is `aria-hidden`. Those are the right instincts. The problem is purely that it runs during the LCP window, doing SVG and clip-path work on the main thread while the browser should be painting content.

CHANGE TO

- Cut the timeline to ~1.8–2.5s total (compress the 1.1s line draw and the 2s shine).
- Start it only after the hero image has decoded, so it never competes with LCP:
`document.querySelector('.cover-img') → new Image().decode().then(startIntro)`
- Optionally show it only on direct/bookmark entries (`!document.referrer`), so search visitors go straight to content.

T8 `/index.html` is a second live copy of `/`, and every internal link points at it **MEDIUM**

NOW

`curl -I https://blakecody.com/index.html` → **200**. The canonical correctly names `https://blakecody.com/`, which resolves the duplicate for Google — but the header brand link and the "Home" nav item both use `href="index.html"`, so every internal link on the site sends crawlers and link equity to the non-canonical variant.

CHANGE TO

```
<a class="brand" href="/">...</a>
<a href="/" class="active">Home</a>
```

Two lines per page, across all five pages.

T9 Sitemap `lastmod` is stale and contradicted by the server **MEDIUM**

NOW

All five URLs claim `<lastmod>2026-06-21</lastmod>`. The live `Last-Modified` header on `/` is **2026-07-20**. The file also uses `<priority>`, which Google has ignored for years.

`lastmod` is the one sitemap field Google does still use, and only while it trusts it. A date that is provably wrong teaches it not to.

CHANGE TO

Generate the sitemap in the same GitHub Action as **T2**, taking each date from the file's last git commit:

```
git log -1 --format=%cs -- index.html # -> 2026-07-20
```

Drop `<priority>`. Keep `404.html` excluded (correct as-is).

T10 Image weight, dimensions and formats MEDIUM

NOW

File	Natural	Displayed	Size	Problem
logo.png	600×600	30×30	45 KB	400× more pixels than used; ~44 KB wasted on every page
cover.jpg	1350×1800	—	443 KB	heaviest asset on the site (used on the Book page)
og-image.png	1200×630	—	91 KB	PNG for a photographic social card; JPEG/WebP ≈ 25 KB
hero-bg.jpg	1600×1000	full-bleed	90 KB	acceptable; WebP ≈ 45 KB
app-icon.png	192×192	132×132	29 KB	fine, slightly oversized

No `<img>` on the page carries `width / height` attributes (confirmed in the render), so each one is a layout-shift opportunity. No `srcset`, and no WebP or AVIF anywhere on the site.

CHANGE TO

- Re-export `logo.png` at 60×60 (2× of its display size) → roughly 4 KB.
- Add explicit `width` and `height` to all three `<img>` elements: ``
- Convert `hero-bg`, `cover` and `app-icon` to WebP, serving JPEG/PNG as a `<picture>` fallback.
- Re-export `og-image` as JPEG at ≤40 KB — social scrapers do not need lossless.

T11 Scripts are not deferred LOW

`<script src="app.js?v=23"></script>` sits at the end of `<body>` — the classic pattern, so the parsing cost is small, but the download does not begin until the parser reaches it. Move both scripts into `<head>` with `defer`: the browser starts fetching immediately and still executes after parse, in order.

```
<script src="app.js?v=23" defer></script>
```

T12 The Substack embed iframe is not lazy LOW

Rendered `loading: "auto"` on `.sub-frame`, whose `src` is set by JS to `https://bccody.substack.com/embed`. That pulls a whole third-party document during load for a signup box that lives below the fold.

```
<iframe class="sub-frame" title="Subscribe by email" loading="lazy" ...>
```

Better still: set `src` from an `IntersectionObserver` when it comes within a viewport of the fold.

T14 Missing feed link, `rel="me"`, and Bing verification LOW

```
<link rel="alternate" type="application/rss+xml"
      title="Blake Cody on Substack" href="https://bccody.substack.com/feed">
```

Add `rel="me"` to the Substack and App Store links so they corroborate the `sameAs` list in **T5**. And verify the site in Bing Webmaster Tools — see **K5**.

T13 Caching is capped at 10 minutes and you cannot change it INFO

Every asset returns `cache-control: max-age=600`. That is GitHub Pages' fixed default — not configurable. Worth knowing that your `?v=23` cache-busting query strings are not buying you anything against a 10-minute TTL.

If repeat-visit performance ever matters, putting Cloudflare's free tier in front lets you set `max-age=31536000` on hashed assets, and hands you HSTS, Brotli and image resizing at the same time.

T15 Already correct — do not touch PASS

- `lang="en"`, `charset` first, viewport meta — all correct.
- Canonical is absolute and self-referential.
- gzip is active on HTML, CSS and JS (6.8 KB HTML → 2.3 KB on the wire).
- `robots.txt` allows all and names the sitemap.
- A missing URL returns a real **404**, not a soft 200.
- Google Search Console verification tag is in place.
- Open Graph and Twitter cards are present with correct 1200×630 dimensions.
- Heading nesting in the static HTML is clean: one H1 → H2 → H3.

B. Content & text

C1 The title tag uses half its budget and targets nothing HIGH

NOW

```
<title>Blake Cody – Aligned & Writing</title> 30 characters
```

Google renders roughly 60 characters. You are using half. "Aligned & Writing" reads as a fragment rather than a claim, and the only searchable term in it is your own name — which you already rank first for and gain nothing by defending.

CHANGE TO

Pick by what you want the homepage to win. Lead with the book:

```
Blake Cody – Author of Aligned: Lead with Love, Live with Intention 66 chars
```

Or lead with the writing:

```
Blake Cody | Essays on Discipline, Love & Intentional Living 59 chars
```

C2 The H1 is just your name HIGH

NOW

```
<h1 class="cover-title">Blake Cody</h1>
<p class="cover-lede">Lead with Love, Live with Intention</p>
```

The single strongest on-page signal spends itself on a brand term. The line that actually describes the site sits in a `<p>`.

CHANGE TO

Swap the semantics without changing the design — the name keeps its visual weight, the H1 carries meaning:

```
<p class="cover-eyebrow">Blake Cody</p>
<h1 class="cover-title">Lead with Love, Live with Intention</h1>
<p class="cover-lede">Essays on discipline, mindset and direction –
  written from inside the work, not after it.</p>
```

If you would rather not touch the layout at all, the one-line version also works: `<h1>Blake Cody – Lead with Love, Live with Intention</h1>`

C3 146 words of static copy — there is nothing for Google to read HIGH

NOW

The complete indexable text of the homepage is the navigation, four section labels, four card blurbs and a footer. No sentence on the page explains who you are, what *Aligned* argues, or what the essays are about. Google has no passage to match any query against except your name.

CHANGE TO

Add a 120–180 word block between the hero and "Latest articles", written for a human first. Work in the language people actually type: *discipline*, *self-discipline*, *mindset*, *intentional living*, *rebuilding after burnout*, *journaling habit*. Add two or three sentences of framing above the "Explore" grid as well. Target roughly 400–500 static words on the page — enough to establish topic without turning the homepage into an essay.

C4 Every article lives on Substack, so this domain has nothing to rank HIGH

NOW

```
<a class="post" href="{p.link}" target="_blank" rel="noopener">
...<span class="post-more">Read on Substack -</span>
```

Your sitemap contains five URLs. None of them is an article. Every essay you publish builds `bccody.substack.com`'s authority, and `blakecody.com` stays a five-page brochure that can only ever rank for "Blake Cody". The three post cards all carry the same non-descriptive call to action.

This is the single highest-value change in the audit, and the only one that changes the ceiling rather than the floor.

CHANGE TO

Host the essays here and let Substack be the distribution channel, not the home:

1. The GitHub Action from **T2** writes each post to `/articles/<slug>.html` with the full text, an `Article` JSON-LD block, and a self-referential canonical.
2. The Substack version links back to the canonical URL on your domain.
3. `articles.html` becomes a real internal hub linking to on-site articles instead of an outbound bounce page.
4. The sitemap grows from 5 URLs to 5 + one per essay, all on your domain.

If you are not ready to move: keep the outbound links, but tag them (**K3**) so you can at least measure what the site sends Substack.

C5 Heading hierarchy breaks the moment posts render MEDIUM

NOW

Rendered DOM order:

```
H2 Latest articles
H2 Why I Started Writing in Public      <- sibling of the section, not a child
H2 The Case for Doing Less, Better
H2 Notes on Building Small Things
H2 Explore
```

CHANGE TO

In `postCard()`, one character: `<h2 class="post-title">` → `<h3 class="post-title">`. This also fixes `articles.html`, which reuses the same function.

C6 The meta description is accurate but gives no reason to click MEDIUM

NOW

```
Blake Cody – author of Aligned: Lead with Love, Live with Intention. Honest
dispatches on love, discipline, and living with intention.           134 chars
```

Good length, but "intention" appears three times and there is no hook or offer.

CHANGE TO

```
Essays on discipline, love and living with intention – from Blake Cody, author
of Aligned. New writing most weeks, plus a free opening chapter.  146 chars
```

The free chapter is a concrete reason to click that you are already offering and not currently advertising in the snippet.

C7 Title, OG title and description all say different things

MEDIUM

NOW

Field	Value
<title>	Blake Cody — Aligned & Writing
og:title	Blake Cody — Honest dispatches from inside the work
meta description	Blake Cody — author of <i>Aligned...</i>
og:description	Rebuilding discipline, mindset, and direction in real time...

Someone who finds you in search and someone who gets your link in a message see two different positioning statements.

CHANGE TO

Choose one line and use it everywhere. Then add the four tags that are missing:

```
<meta property="og:locale" content="en_US" />
<meta property="og:image:alt" content="Blake Cody – Lead with Love, Live with Intention" />
<meta name="twitter:image:alt" content="Blake Cody – Lead with Love, Live with Intention" />
<meta name="twitter:creator" content="@yourhandle" /> if you have one
```

C8 Post thumbnails ship with empty alt text

MEDIUM

```
// app.js, postCard()
`<div class="post-thumb"></div>`
```

An empty `alt` declares an image decorative. These thumbnails are the post's cover art — they carry meaning, and they are the only images on the page that could earn image-search traffic.

```
``
```

(The header logo's `alt=""` is correct — the words "Blake Cody" sit right next to it.)

C10 No author signals on the page

MEDIUM

For a personal-brand site giving advice on discipline and mindset, Google leans heavily on author and entity signals. The homepage has no byline, no bio text, no visible dates and no links to profiles — nothing that says a real, identifiable person stands behind the advice.

CHANGE TO

Add a compact author block above the footer: photo, two sentences in the first person, and links to Substack, the App Store and any social profiles. Back it with the `Person` node from **T5** and mark the links `rel="me"` (**T14**). Add `dateModified` to the `WebPage` node so freshness is machine-readable.

C11 No keyword target has been chosen for this page MEDIUM

Nothing on the site currently targets a query other than "blake cody". That is a strategy gap, not a page defect — but it means the homepage is being asked to do a job it is the wrong page for.

CHANGE TO

Page	Should target
Homepage	Brand + book: <i>Blake Cody, Aligned book, lead with love live with intention</i> . Optimize it as an entity hub, not a topic page.
/articles + article pages	The topics: <i>how to rebuild discipline, living with intention, daily journaling habit</i> . Requires C4.
/book	<i>Aligned Blake Cody, free first chapter</i>
/reflect	<i>free reflection prompts, write to a stranger</i>

Stop asking the homepage to rank for topics. That is what the articles are for — which is why **C4** is the change that matters most.

C9 Already correct — do not touch PASS

- Explore card anchors are descriptive: "Read all →", "Open the book →", "Start writing →", "More about me →".
- App Store badge has a real `alt`: "Download My JRNL on the App Store".
- App icon has a real `alt`.
- The scroll cue has an `aria-label` despite having no text.
- The header logo correctly uses `alt=""` next to its text equivalent.
- Copy voice is consistent and genuinely well written — the writing is not the problem, the amount of it is.

C. Tracking & measurement

K1 GoatCounter counts pageviews and nothing else

HIGH

NOW

```
<script data-goatcounter="https://bccody03.goatcounter.com/count" async
  src="//gc.zgo.at/count.js"></script>
```

The default snippet, unmodified. A grep of `app.js` for `goatcounter` returns zero hits. You can see that the homepage got traffic. You cannot see whether a single person clicked through to Substack, tapped the App Store badge, requested the chapter, or sent a reflection. **Every conversion on this site is currently unmeasured**, which means you have no way to tell whether any of the other changes in this document worked.

CHANGE TO

Add a helper and five call sites. GoatCounter events are free and need no extra script:

```
function track(name, title) {
  if (window.goatcounter && window.goatcounter.count) {
    window.goatcounter.count({ path: name, title: title || name, event: true });
  }
}

// 1 + 2 – every outbound click, one delegated listener
document.addEventListener('click', (e) => {
  const a = e.target.closest('a'); if (!a || !a.href) return;
  if (a.href.includes('substack.com'))
    track('click-substack-' + (a.id || a.closest('.post') ? 'card' : 'other'), 'Substack click');
  if (a.href.includes('apps.apple.com')) track('click-appstore', 'App Store click');
});

// 3 – in the chapter form's success branch (app.js ~line 240)
track('chapter-request', 'Chapter 1 requested');

// 4 – in the reflect form's success branch (app.js ~line 305)
track('reflect-submit-' + (payload.category || 'none'), 'Reflection sent');

// 5 – in the init() catch block – turns T3 into an alert, not a surprise
track('feed-error', 'Substack feed failed');
```

Those five rows in GoatCounter are the entire business value of the site, and right now none of them exists.

K3 Outbound links carry no attribution, so Substack cannot credit you HIGH

NOW

Every Substack link — nav, footer, book CTA, all three post cards — points at the bare URL. In Substack's own stats they all land as direct/unknown traffic. The number you most want ("how many subscribers did my website send me?") is currently unknowable from either end.

CHANGE TO

```
// app.js — where substackLinks are pointed at CONFIG.substackUrl
const UTM = '?utm_source=blakecody.com&utm_medium=referral&utm_campaign=site';
els.substackLinks.forEach((a) => {
  if (!a) return;
  a.href = CONFIG.substackUrl + UTM + '&utm_content=' + (a.id || 'link');
});

// postCard() — mind the & vs ? if the post URL already has a query
const sep = p.link.includes('?') ? '&' : '?';
const href = p.link + sep + 'utm_source=blakecody.com&utm_medium=referral'
  + '&utm_campaign=home-latest';
```

Use a distinct `utm_content` per placement — `nav`, `footer`, `book-cta`, `card` — so you learn which placement actually converts rather than which one gets clicked.

K4 Search Console is verified but barely used MEDIUM

The `google-site-verification` meta tag is present, so the property exists. Checklist to make it useful:

- Add a **Domain property** (DNS-verified) alongside the URL-prefix one. A domain property covers http, https and www in a single view — which matters a great deal given **T1**.
- **Submit `sitemap.xml` explicitly.** robots.txt discovery works but is slower and gives you no error reporting.
- Run **URL Inspection** → **Test live URL** → **View rendered HTML** on `/`. Confirm the three post cards are present and are *real posts*, not the samples from **T3**. This single test settles both **T2** and **T3** definitively.
- Watch the Coverage report for `/index.html` appearing as a duplicate (**T8**).

K5 Not verified in Bing Webmaster Tools MEDIUM

Two clicks: `bing.com/webmasters` → **Import from Google Search Console**. No new tag needed.

Worth doing beyond Bing's own traffic: Bing's index feeds Copilot and parts of ChatGPT's search. For a personal brand whose goal is to be the answer to "who is Blake Cody", being indexed by the engine those assistants read is not optional — and Bing does not execute JavaScript, which makes **T2** its blocker.

K6 Form submissions cannot be traced to a source MEDIUM

NOW

```
body: JSON.stringify({ email, request: 'Chapter 1 PDF', submittedAt })
body: JSON.stringify({ category, message, name, email, submittedAt })
```

Both Formspree payloads record what was sent but nothing about where it came from. "Did the chapter downloads come from search or from the newsletter?" is currently unanswerable.

CHANGE TO

```
// capture the landing context once, on first page of the session
try {
  if (!sessionStorage.getItem('entry')) sessionStorage.setItem('entry', JSON.stringify({
    landing: location.pathname, referrer: document.referrer, query: location.search
  }));
} catch (e) {}

// then include it in both payloads
entry: sessionStorage.getItem('entry'), page: location.pathname
```

K7 No visibility when something breaks MEDIUM

If rss2json goes down, or the API key hits its quota, or a GitHub Pages certificate renewal fails, you find out by visiting your own site. Two free fixes:

- The `feed-error` event from **K1** — a GoatCounter row that spikes the moment the feed breaks.
- An UptimeRobot check (free tier) on <https://blakecody.com>, 5-minute interval, email alert. It catches cert failures and DNS problems too.

K2 Protocol-relative analytics URL LOW

```
src="//gc.zgo.at/count.js" → src="https://gc.zgo.at/count.js"
```

A 2013 pattern. Combined with **T1**, on an `http://` load it currently fetches your analytics over plaintext.

K8 No Core Web Vitals field data LOW

Every performance number in this document is a lab measurement. You have no data on what real visitors experience.

- **Free and immediate:** run PageSpeed Insights on `/` monthly and read the CrUX panel (it needs enough traffic to populate).
- **Real field data:** the `web-vitals` library is 1.5 KB — bucket the values and push them into GoatCounter as events: `track('cwv-lcp-' + bucket)`. That gives you distributions for LCP, INP and CLS at zero cost and with no extra vendor.

K9 No privacy page LOW

GoatCounter is cookieless and IP-anonymising, so you are genuinely clear of a consent banner — a good choice, worth keeping. But Formspree stores names, emails and personal reflections, and there is no page saying what happens to them.

Add a short `/privacy.html` : what GoatCounter records, what the two forms store, how to ask for deletion. Link it in the footer. It is a trust signal Google's quality systems look for, and it costs twenty minutes.

Evidence log

Check	Result
<code>curl -I https://blakecody.com</code>	200 · GitHub.com via Fastly · Last-Modified: Mon, 20 Jul 2026 · gzip · max-age=600
<code>curl -I http://blakecody.com</code>	200, no redirect → T1
<code>curl -I .../index.html</code>	200 — duplicate of / → T8
<code>curl -I .../does-not-exist</code>	404 — correct
DNS	apex → 185.199.108–111.153 · www → bccody03.github.io (www redirect not verifiable from this sandbox)
Static text extraction	146 words in <body> → C3
Headless render (Chromium 1194)	242 words · 3 post cards · 0 JSON-LD blocks · FCP 12.9 s with fonts unreachable · coverBg set by JS · no width / height on any image
Feed fallback	"Showing sample posts" + 3 fabricated posts with href="#" → T3
Third-party runtime deps	fonts.googleapis.com (blocking) · api.rss2json.com (content) · gc.zgo.at (analytics) · bccody.substack.com (iframe)
Asset weights	cover.jpg 443 KB · og-image.png 91 KB · hero-bg.jpg 90 KB · logo.png 45 KB at 30×30 → T10
Not verifiable here	api.rss2json.com and www.blakecody.com are blocked by this sandbox's egress policy — both flagged for manual confirmation

Page 01 of 6 complete. Remaining: `articles.html` · `book.html` · `reflect.html` · `about.html` · `404.html`.

Findings marked T2, C4, C5, C8, K1 and K3 live in `app.js`, which is shared across every page — fixing them once fixes them site-wide. They will be cross-referenced rather than repeated in later reports.

blakecody.com SEO change log · generated 3 September 2026