

Page 07 — Privacy /privacy.html

Audited: 7 September 2026 · repository `bccody03/blog-site` @ `main c0a998a`
Live URL: `https://blakecody.com/privacy.html` — HTTP 200
Method: static HTML parse · heading-tree extraction · **claim-by-claim verification of the policy text against the actual implementation in `app.js`**
Status: Audit only — **no code changes applied**. This page was created during SEO pass 1, in response to finding **K9** on the homepage.

4 CHANGES PROPOSED	0 HIGH	2 MEDIUM	2 LOW
-----------------------	-----------	-------------	----------

Two things about this page are better than the rest of the site VERIFIED

1. It is the only page with a correct heading tree.

```
H1 Privacy
H2 Visits
H2 The two forms
H2 Newsletter
H2 Asking for deletion
```

No skipped level, no reversal, every heading describing real content. Every other page audited has either an H1 → H3 jump, an H2 arriving after H3s, or both. Whoever wrote this one got it right without being asked.

2. The policy is true. I checked each claim against the code rather than taking it at face value, because privacy policies are usually written from a template and then quietly drift from what the site does:

Claim	Verified against
"GoatCounter ... sets no cookies, doesn't store your IP address"	✓ accurate — GoatCounter is cookieless by design and does not retain IPs. The conclusion drawn from it — "which is why there's no cookie banner here" — follows correctly.
"counts page views and a handful of clicks (for example, how many people tap through to Substack or the App Store)"	✓ exactly what the delegated listener in <code>app.js</code> does
"Both forms ... are handled by Formspree"	✓ two endpoints, <code>coachWebhook</code> and <code>chapterWebhook</code>
"Reflect ... stores your message, the topic you picked, your email address, and your first name if you gave one"	✓ matches the payload field for field
"The free chapter stores your email address ... doesn't subscribe you to anything"	✓ separate endpoint, no newsletter call in that path
"Each submission also records which page it came from and how you arrived at the site"	✓ this is <code>entryContext()</code> — landing page, referrer and query — spread into both payloads
"The email signup box is Substack's"	✓ a cross-origin iframe to <code>/embed</code>

Every claim holds. That is rare enough to be worth stating plainly — and it is the reason the four findings below are all minor.

A. Content

C39 The deletion route runs through the form that creates the records

MEDIUM

NOW

*"...reply to any email you've received from me — or send a note through the **Reflect form** asking me to delete your previous messages..."*

Two problems, one practical and one structural.

Practical: submitting the Reflect form to request deletion creates a new Formspree record containing a message, an email address and the entry context — so asking to be forgotten generates one more thing to forget. It works, but a careful reader notices, and this is a page read by careful readers.

Structural: "reply to any email you've received from me" only works for someone who has already received one. A person who submitted the chapter form, never got the PDF, and wants their address removed has no route that does not involve submitting another form.

There is no plain email address anywhere on the site — I checked every page.

FIX

Publish a contact address as the primary route and keep the form as the fallback:

```
<p>Email me at <a href="mailto:hello@blakecody.com">hello@blakecody.com</a>
and I'll delete anything you've sent. If you'd rather not email, the
<a href="reflect.html">Reflect form</a> reaches me too.</p>
```

A contactable address is also a small trust and entity signal in its own right, and it can go in the `Person` schema as `email` once it exists.

To be clear about the limits of this: I am not a lawyer and this is not legal advice. I am flagging it as a practical gap a reader will notice, not as a compliance ruling. If the site ever takes payment or serves the EU or UK at scale, that is the point to get an actual opinion.

T39 "Last updated September 2026" is prose only

MEDIUM

NOW

```
<p class="reflect-note">Last updated September 2026.</p>
```

On a privacy policy the date is not decoration — it is how a reader decides whether the document describes current practice. Right now it is an untagged sentence in a class named after a different page's component, with no machine-readable equivalent and no day.

FIX

```
<p class="policy-updated">Last updated  
<time datetime="2026-09-04">4 September 2026</time>.</p>
```

And declare it in schema, which also closes **T40** below:

```
<script type="application/ld+json">  
{ "@context": "https://schema.org",  
  "@graph": [  
    { "@type": "WebPage",  
      "@id": "https://blakecody.com/privacy.html#webpage",  
      "url": "https://blakecody.com/privacy.html",  
      "name": "Privacy",  
      "dateModified": "2026-09-04",  
      "isPartOf": { "@id": "https://blakecody.com/#website" } },  
    { "@type": "BreadcrumbList",  
      "itemListElement": [  
        { "@type": "ListItem", "position": 1, "name": "Home",  
          "item": "https://blakecody.com/" },  
        { "@type": "ListItem", "position": 2, "name": "Privacy" } ] } ] }  
</script>
```

Add a line to whatever checklist you keep: **if the tracking or the forms change, this date changes too.** A policy dated correctly but describing old behaviour is worse than an undated one.

B. Technical

T40 No structured data

LOW

0 JSON-LD blocks — the same gap as `book.html` (**T23**), `reflect.html` (**T27**) and `about.html` (**T32**), and the lowest-priority instance of it on the site.

The snippet in **T39** covers it. Do this one last, or fold it in when you do the others.

C40 "Privacy" is the right title and the right H1 — leave them KEEP

Measured: title 20 chars, H1 one word. On every other page in this audit I have argued that a one-word H1 wastes the site's strongest on-page signal — **C2**, **C15**, **C21**, **C28**, **C34**.

Here the opposite applies, and I want to say so explicitly so it does not get "fixed" for consistency.

- Nobody arrives at a privacy policy from search wanting a compelling headline. They arrive from the footer, looking for one specific word, and finding it immediately is the whole job.
- There is no search upside to dress up — you do not want to rank for anything here, and ranking for "privacy" would be neither achievable nor useful.
- Plain naming is itself a trust signal on this kind of page. "Privacy" reads as straightforward; anything more crafted reads as marketing, on the one page where marketing voice actively costs you credibility.

The `archive-sub` line already does the persuasive work and does it well — *"Short version: this site keeps almost nothing about you, and what it does keep, you can ask me to delete."* That is the best sentence on the page and it is in exactly the right place.

T41 Correct as it stands — do not change PASS

- **Indexable and in the sitemap** ✓ — correct. Some sites `noindex` their policy; a visible, indexable privacy page is a positive quality and trust signal, and Google expects a site collecting emails to have one.
- **Canonical, self-referential** ✓
- **Links out to each processor's own policy** — GoatCounter, Formspree, Substack ✓. Naming your processors is more than most personal sites do.
- **Plain English, no boilerplate.** No "we may share with third parties for legitimate business purposes". It reads like a person explaining what happens, which is both better writing and more legally meaningful than copied template text.
- **267 static words** — the second-highest on the site after the book page, and every one of them load-bearing.
- Footer link present on every page ✓ (**C31** on the Reflect page asks for one *additional* link at the point of disclosure — that is about placement, not about this page's content).

Evidence log

Check	Result
Static text	267 words — second-highest on the site
Title / description	20 chars / 112 chars — description accurate and specific
Headings	H1 → H2 ×4 — the only clean heading tree on the site
Policy accuracy	7 claims checked against <code>app.js</code> ; 7 hold
Canonical / sitemap / robots	self-referential ✓ · included ✓ · indexable ✓ — all correct

JSON-LD	0 blocks
Last-updated date	prose only, month precision, no <code><time></code> , no <code>dateModified</code>
Contact route	No email address anywhere on the site; deletion route runs through the Reflect form
Processor links	goatcounter.com, formspree.io, substack.com/privacy — all present

Page 07 of 7 — the SEO audit is complete. Every page on blakecody.com now has a change log: [index](#) · [articles](#) · [book](#) · [reflect](#) · [about](#) · [404](#) · [privacy](#).

See [seo/PROGRESS.md](#) for the consolidated status of every finding across all seven passes — what is fixed, what is open, and what is still unverifiable from the audit environment.

The UI/visual design track remains at page 01 of 7.

blakecody.com SEO change log · generated 7 September 2026